



AUTARCH | THE GOVERNED AGENT FACTORY

Governed AI agents for finance, audit, and taxation

Every conclusion bounded, approved, and provable.

Enterprise pitch | September 2026

Finance wants AI. It cannot compromise control.

AI now reviews ledgers, tests controls, interprets tax, and prepares decisions.

- Broad ERP access and prompt-level guardrails
- Inconsistent controls rebuilt for every finance pilot
- Activity logs without workpaper-grade evidence
- Approvals disconnected from deployed configuration
- Material findings that still require accountable judgment

THE REAL BOTTLENECK

Financial control, evidence, and accountability at scale.

The category: a governed agent factory

One finance control plane. Many governed specialist agents.

CONVENTIONAL BUILDER


- Prompts + tools
- Model compliance
- Project-specific controls
- Logs after the event



AUTARCH

- Authority + policy + proof
- Rules outside the model
- Reusable DomainPacks
- Evidence by construction

From requirement to governed deployment



The factory creates declarative configuration - never arbitrary executable code.

AI may propose the agent. The governance plane decides whether it exists.

Universal governance. Finance-specific trust.



SHARED PLATFORM


Orchestration and memory
Capability security and budgets
Approvals and lifecycle
Static proof and signed evidence

FINANCE DOMAINPACKS

Accounting, audit + tax workflows
Materiality + approval thresholds
ERP and evidence boundaries
Jurisdiction rules + evaluations

Scale across the office of the CFO without weakening financial controls.


The deterministic governance stack



1	RBAC	Who may wield authority?
2	CAPABILITY KERNEL	What exact action?
3	SCOPE + LIMITS	Where and how much?
4	POLICY-AS-CODE	Deny or escalate?
5	ECONOMIC KERNEL	Within budget and risk?
6	APPROVAL PLANE	Who ratified it?
7	SIGNED LEDGER	Can we prove it?

The kernel cannot be prompted out of policy.

Intelligence above. Deterministic control below.



No planner, framework, model, or tool bypasses the chokepoint.

One governed finance action: end-to-end sequence

01 INTENT
Finance request + context

03 NORMALIZE
Typed adapter parameters

05 EVALUATE
Policy + budget

07 EXECUTE
Trusted adapter

09 SEAL
Sign rationale + outcome

02 DELIBERATE
Propose + challenge

04 AUTHORIZE
Grant + scope + limits

06 PRESIDE
Rule or quorum approval

08 ASSESS
Quality + safety panel

10 EMIT
Events + compliance evidence

A failed gate prevents execution - and the refusal remains explainable.

Security architecture: authority is structural



IDENTITY	Principal + node identity	Attributable actor/runtime
ENTITLEMENT	RBAC-filtered grants	Role-aligned access
LEAST PRIVILEGE	Exact/wildcard capabilities	No ambient authority
CONFINEMENT	Scopes + quantitative limits	Bounded records/amounts
DELEGATION	Monotonic attenuation	Child never exceeds parent
INTEGRITY	Ed25519 + hash chain	Tamper-evident evidence

Prompt injection may alter a proposal. It cannot manufacture a grant.

Memory, state, and evidence are different planes

WORKING CONTEXT

Current reasoning
Council + run state

GOVERNED RECALL

Semantic + episodic
Long-term knowledge

PRECEDENT

Prior human rulings
Reusable decisions

DURABILITY

Resume safely
Avoid double execution

WHY-MEMORY


Signed rationale
Outcome + proof

OPERATIONS

Events + telemetry
Monitoring + export

Conversation history is useful context. It is not, by itself, audit evidence.

Portable control from finance UI to system of record



Laptop | container | air-gapped host | distributed mesh

Approval is bound to what actually ships

SHA-256

CONTENT FINGERPRINT

NO STALE APPROVAL

NO SILENT PRIVILEGE EXPANSION

- Semantic versions become immutable after registration
- Quorum approval binds pack + blueprint + fingerprint
- Any change invalidates the prior approval
- RBAC narrows authority again at compilation
- Empty grants mean no authority by default

Prove safety before execution

FORBID

file.delete can never execute

REQUIRE APPROVAL

payment.send never auto-executes


CONFINE

file.write stays inside reports/

- Proof applies to every attenuated child agent
- Failure returns a concrete counterexample
- Compilation fails closed when an invariant does not hold

Not a model promising good behavior - a control plane proving a boundary.

Multi-agent work - without authority sprawl



Least privilege | isolated tools | parallel waves | unified audit chain

Factory internals: safe specialization at scale

- 1 Finance requirement
- 2 Select Finance DomainPack
- 3 Resolve versioned AgentBlueprint
- 4 Validate tools + grants + policy + budget
- 5 **Prove mandatory invariants**
- 6 **Bind quorum approval to fingerprint**
- 7 Filter through principal RBAC
- 8 Compile, monitor, retire deployment

Scale agents without scaling unreviewed privilege.

Evidence is part of execution

WHY

Intent
Proposal + critique
Recommendation

CONTROL

Kernel decision
Policy + budget
Human approval


PROOF

Tool outcome
Evaluation verdict
Identity + signature

From agent-asserted compliance to verifiable operational evidence.

Retention controls can redact sensitive fields while preserving chain integrity.

Continuous audit and tax assurance



Ledgers | controls | accounting treatment | tax positions | signed evidence

The reusable pattern

ingest > reason > verify > approve > act > prove

Finance control pattern - not autonomous audit opinion or tax advice.

One platform across the office of the CFO



INTERNAL AUDIT

Control testing | sampling | evidence | findings

TAXATION

Indirect tax | provisions | jurisdiction monitoring

CONTROLLERSHIP

Journal review | reconciliations | close

TREASURY + RISK

Cash forecasting | payment controls | exposure

FP&A

Variance analysis | scenarios | management reporting

Land one finance process. Reuse the controls across the CFO portfolio.

Fit into the stack customers already own

TOOLS

MCP
LangChain
Microsoft Agent Framework

FINANCE DATA

ERP + consolidation
Tax + treasury systems
Ledgers + workpapers


MODELS

Azure OpenAI + OpenAI
Anthropic
Ollama / local

Autarch is the governance substrate beneath the ecosystem - not another model silo.

Deploy from edge to enterprise

- Pure Python core + SQLite; zero required runtime dependencies
- Offline path and first-class local-model support
- Container health/readiness and structured telemetry
- Durable runs, retries, rate limits, and circuit breakers
- Encrypted mesh and signed distributed provenance
- Typed errors for predictable integration



Business value

SHIP FASTER

Reuse blueprints, controls, integrations, and approval patterns.

ACCELERATE ASSURANCE

Generate operational evidence as part of execution - not after it.

REDUCE BLAST RADIUS

Give every agent only the tools, scope, limits, and budget it needs.

AVOID LOCK-IN

Replace models and orchestration while governance remains consistent.

Competitive position

CAPABILITY	AUTARCH	TYPICAL ORCHESTRATION-FIRST STACK
Versioned agent factory	NATIVE	Application-built
Deny-by-default capability kernel	NATIVE	Tool / auth layer
Static safety guarantees	NATIVE	Tests / policy checks
Content-bound quorum approval	NATIVE	Workflow integration
Signed action + evaluation evidence	NATIVE	Logging integration
Model and tool ecosystems	INTEGRATES	Often mature / native

Complement orchestration frameworks. Own the governance layer.

The market has runtimes. Autarch adds a control plane.



MICROSOFT AGENT FRAMEWORK	Typed agents + graph workflows	Microsoft ecosystem, middleware, sessions
LANGGRAPH	Low-level stateful orchestration	Durable graphs, interrupts, LangSmith
CREWAI	Role-based crews + event-driven flows	Accessible teams, state, persistence
OPENAI AGENTS SDK	Lightweight agents + handoffs	Minimal primitives, tools, tracing
AUTOGEN	Event-driven multi-agent runtime	Messaging and team patterns
AUTARCH	Governed creation + consequence	Authority, proof, signed evidence

Overlap is real. Architectural optimization is different.

Autarch vs Microsoft Agent Framework (MAF)

MAF STRENGTHS

- Typed agents + sessions
- Middleware + provider integrations
- Graph workflows + checkpointing
- Streaming + human request/response
- Microsoft-supported ecosystem

AUTARCH DIFFERENTIATION

- Scoped capability authority
- Monotonic delegation
- Pre-run static safety proof
- Fingerprint-bound quorum approval
- Signed rationale + outcome evidence

Winning combination: MAF orchestrates. Autarch governs consequential functions.

Autarch vs LangGraph

LANGGRAPH STRENGTHS

- Deterministic + agentic graph nodes
- Durable checkpoints + stores
- Interrupts + human review
- Time travel + state inspection
- LangSmith deployment + observability

AUTARCH DIFFERENTIATION

- Kernel primitive authorization
- Actor/quorum-aware approval
- Uniform side-effect gating
- Safety proof before graph execution
- Cryptographically attributable evidence

Winning combination: keep LangGraph state machines; govern actions behind tool nodes.

Autarch vs CrewAI and AutoGen

CREWS / TEAMS STRENGTHS

- Intuitive roles + delegation
- Collaborative agent patterns
- State + memory + persistence
- Flexible flow coordination
- Fast multi-agent prototyping

AUTARCH DIFFERENTIATION

- Delegation narrows real authority
- Tool possession is not permission
- Policy + economic kernels
- Fleet-wide attenuation guarantees
- Signed accepted + blocked actions

Winning combination: crews reason and collaborate; Autarch governs authority and evidence.

Autarch vs OpenAI Agents SDK

OPENAI SDK STRENGTHS

- Productive minimal primitives
- Tools, handoffs + sessions
- Input/output/tool guardrails
- Human-in-the-loop runs
- Deep OpenAI + realtime integration

AUTARCH DIFFERENTIATION

- Provider-neutral control contract
- Grant/scope/limit authorization
- Sound unconditional-policy proof
- Durable quorum + content binding
- Locally verifiable signed evidence

Winning combination: OpenAI agents reason and hand off; Autarch controls side effects.

Where Autarch supersedes - and where it complements

CAPABILITY	AUTARCH	MAF	LANGGRAPH	CREWAI	OPENAI
Agent / workflow orchestration	YES	LEAD	LEAD	LEAD	YES
Durable state / resume	YES	LEAD	LEAD	YES	YES
Deny-by-default action authority	LEAD	CUSTOM	CUSTOM	CUSTOM	GUARD
Scoped + limited delegation	LEAD	CUSTOM	CUSTOM	CUSTOM	CUSTOM
Static safety proof	LEAD	-	-	-	-
Content-bound quorum approval	LEAD	CUSTOM	CUSTOM	CUSTOM	CUSTOM
Signed tamper-evident action proof	LEAD	TRACE	TRACE	TRACE	TRACE

Supersedes at governance. Complements at orchestration and ecosystem.

The winning architecture is compositional

MAF | LANGGRAPH | CREWAI | OPENAI AGENTS

planning | state | collaboration | developer experience

V

AUTARCH GOVERNANCE PLANE

identity | authority | policy | budget | approval | proof

V

ERP | TAX | TREASURY | FILES | APIs | MCP TOOLS

Become the layer no consequential agent should bypass.

Go-to-market: land with one governed workflow

01

FINANCE PARTNER

Select an evidence-heavy process with a control owner.

02

FINANCE PACK

Encode accounting, tax, authority, and evaluation.

03

CONTROLLED PILOT

Measure coverage, reviewer hours, and false positives.

04

CFO EXPANSION

Expand across audit, tax, close, treasury, and FP&A.

Commercial wedge: governed finance workflows + reusable control packs.

Readiness and responsible roadmap



AVAILABLE NOW

Governed agents + orchestration
Capabilities, policy, RBAC, budgets
Approvals + static guarantees
Signed provenance + evaluation
Blueprints, DomainPacks + lifecycle

NEXT VALIDATION MILESTONES

External security review
Design-partner outcomes
Signed DomainPack distribution
Enterprise registry service
Governed model-assisted proposals

Production-oriented architecture. Production claims earned through validation.

Competitive research sources

MICROSOFT AGENT FRAMEWORK

learn.microsoft.com/agent-framework/

LANGGRAPH

docs.langchain.com/oss/python/langgraph/overview

CREWAI

docs.crewai.com/

AUTOGEN

microsoft.github.io/autogen/stable/

OPENAI AGENTS SDK

openai.github.io/openai-agents-python/

Reviewed 2 September 2026. Comparisons describe documented architectural emphasis; revalidate product versions, deployment requirements, and licensing during diligence.

THE ENTERPRISE AGENT CONTROL PLANE

Scale finance intelligence.

Preserve financial control.

AUTARCH - THE GOVERNED AGENT FACTORY

From finance requirement to approved, bounded, and provable deployment.

AUTARCH

You don't use AI. You preside over it.

Finance-controlled | Audit-ready | Model-agnostic | Provable

```
pip install autarch
```